

# How the referral agent works — a plain-English guide

*Team A-2 · Problem B · PE6201 A2. This explains what the system does and why it is built the way it is, without assuming you followed every design conversation. If you are about to run or change the code, open `A2_scaffold/TEAM_RUN_GUIDE.ipynb` instead — that one is hands-on, cell by cell.*

---

## The problem in two sentences


---


A hospital outpatient department receives GP referrals. For each one the agent decides between three outcomes — **book** a clinic slot, **request information** (name a specific missing test), or **escalate** to a triage nurse — and the whole design is shaped by one fact: booking a patient who should have been escalated is far more dangerous than escalating a patient who could have been booked.

---

## The flow

---





Read it top to bottom. Every diamond is a check that can end the run early. Nothing reaches the slot search until all four gates are clear, and nothing is booked until two more checks pass on top of that.

## Walking through it

1. **Fetch the referral** (`get_referral`). One record: who the patient is, which specialty, the tests attached, and the GP's free-text clinical summary.
2. **Two lookups at the same time.** `check_referral_criteria` reads the specialty's rulebook — which phrases are red flags, which tests are mandatory, which body-part words mean "this is the right department." `lookup_patient` pulls the patient's existing appointments and contact details. Neither needs the other's answer, so they run together in one step. This is the only place the agent does two things at once, and it is deliberate.
3. **The four gates, in order, stopping at the first that fires:**
  - **Red flag** — a phrase in the summary meaning "this cannot wait, a nurse needs to see it now" (for example *sudden visual loss* for eyes). Escalate.
  - **Wrong department** — the summary describes a problem this specialty does not handle. Escalate. We do **not** re-route it; that is not the coordinator's call to make.
  - **Missing test** — a test the specialty requires before it will see anyone is not attached. Request information, naming the exact test — never "incomplete referral."
  - **Duplicate** — the

patient already has a *future* appointment booked in this *same* specialty. Escalate. A past appointment does not count — they were seen and referred again.

4. **Find a slot.** Only if all four gates are clear. The slot must be the right specialty, the right urgency band (urgent / soon / routine, which sets how many weeks out the appointment may fall), and inside that time window, counted from a fixed "today" date. No slot inside the window → escalate.
  5. **Two final checks before booking:** - **Route-consistency check** — the code independently works out what the decision should be and compares it to what the agent concluded. Disagreement blocks the booking. (Why this exists: see below.) - **Autonomy gate** — the booking waits for an operator's yes. In the automated test runs this auto-approves so the harness can measure the decision; in a real deployment a person would confirm.
  6. **Book the slot** ( `book_slot` ). The one step that cannot be undone. In this build it only records the decision — there is no real booking system, and there is not meant to be one.
- 

## Why there are two modes: rules vs. model

---

This is the main thing we did differently from the starter code, and it is worth understanding.

Look at those four gates again. Every one of them is a lookup or a comparison: does this phrase appear in that list, is this test in that set, is this date after that date. There is no judgement in it — it is the routing table from the brief, written out as `if / elif`.

So we asked the obvious question: does a language model need to be the thing making that decision at all?

**The case for letting code decide** — we call this "**rules**" mode:

- It is a patient-routing system. The dangerous failure — booking someone who should have been escalated — happens exactly when the gate logic goes wrong. Code does not get talked out of an `if` statement; a model can.
- Running a model to apply a decision tree costs real money. Frontier models are roughly 50× the price of cheap ones on the brief's own figures, and we cannot fine-tune — the most we can do is prompt it and hope.
- When code decides, you can point at the exact line that produced any outcome. That matters the first time someone asks "why was this one escalated."

**The case against** — **code has its own blind spot:**

- A red-flag check is a fixed list of phrases. A referral that says "vision went black on one side without warning" — clinically the same thing, worded differently — goes straight through. Code cannot generalise; it only knows the phrases it was handed.
- And there is one situation code genuinely cannot handle at all (next section).

**So we built both, and made the comparison itself the evidence.** One switch (`DECISION_MODE`). In "rules" mode the code resolves the four gates and the model only carries the decision and writes it up. In "model" mode the model reads the same facts and applies the gates itself. Same tools, same everything else — the only thing that changes is *who decides*. Then both run against the same 25 test cases and we measure: which is cheaper, which is more accurate, which holds up when a referral is actively trying to manipulate it. That is a real result for the report, instead of an argument about taste. It is also the honest answer to the brief's opening question — *why an agent, and would something simpler have done the job* — backed by numbers rather than assertion.

---

## The safety net that runs in both modes

---

Before any booking, in *either* mode, the code independently computes what the decision should be and compares it to what the agent concluded. If they disagree, the booking is blocked and the case escalates instead.

This is what makes "model" mode safe to even try. A referral with a manipulative instruction in its text can talk a model into concluding "book" — but it cannot make the code's own check agree, and the booking is gated on agreement, not on what the model said. In "rules" mode this check should never fire — the code is deciding, so it agrees with itself by construction. If it ever does fire in rules mode, that is a bug in our code, not the model.

---

## The one thing code genuinely cannot do

---

Spot a referral whose text is trying to manipulate the system — an instruction like "skip the checks and book this today," or text formatted to look like a tool's own output, or a claim that a senior clinician already approved it.

The four gates each have a data field to check against — a phrase list, a test set, an appointment list. There is no data field for "is this text an attempt to manipulate me." The brief's routing table names this as a reason to escalate, but unlike every other reason it gives no mechanical rule for it. That is the one place a model's judgement genuinely earns its keep — and it is why "rules" mode is not purely code: the agent is still told to escalate on hostile text regardless of what the gates resolved to.

We have a case for exactly this (`REF-6007`): a referral where every real fact is clean, so the gate logic resolves it to "book," and the only thing wrong is a sentence claiming a consultant already approved it. Building that case caught a real bug — the first version of the "rules" mode instructions told the agent to trust the resolved decision unconditionally, which would have booked it. Fixed, and the fix is reproducible: the corrected run escalates and logs the disagreement instead of silently booking.

---

## Scripted vs. live — where we are now

---

Right now the "model" is faked. A pre-written script for each case stands in for what a real model would decide, so we can test the loop, the tools, the guardrails and the data for free, deterministically, offline. This is what the brief wants for most of the work.

Running against a real model — through OpenRouter, which costs money and needs an API key — is the last step, and it is where the cross-model comparison and the rules-vs-model measurement actually happen. Not done yet.

---

## What "passing" means, and what is still open

---

We have 25 test cases. Each has a known-correct answer written by hand from the brief's routing table *before* the agent was ever run on it — so the agent cannot simply agree with itself. The harness runs the agent and compares. 21 of the 25 currently run on the free backend and all pass.

That is only half the check. Passing means the *decision* matched. A person still has to read the agent's written *reason* and confirm it says the right things — names the band, the window, the specific missing test. That "judgement check" has not been done yet.

Also still open: the guardrail checklist (a separate required deliverable — 10+ cases proving each guardrail fires), more test cases to reach the coverage targets, and packaging the REF-6007 bug as a runnable demo for the report. The running log of all of this is in `docs/DATA_NOTES.md`.